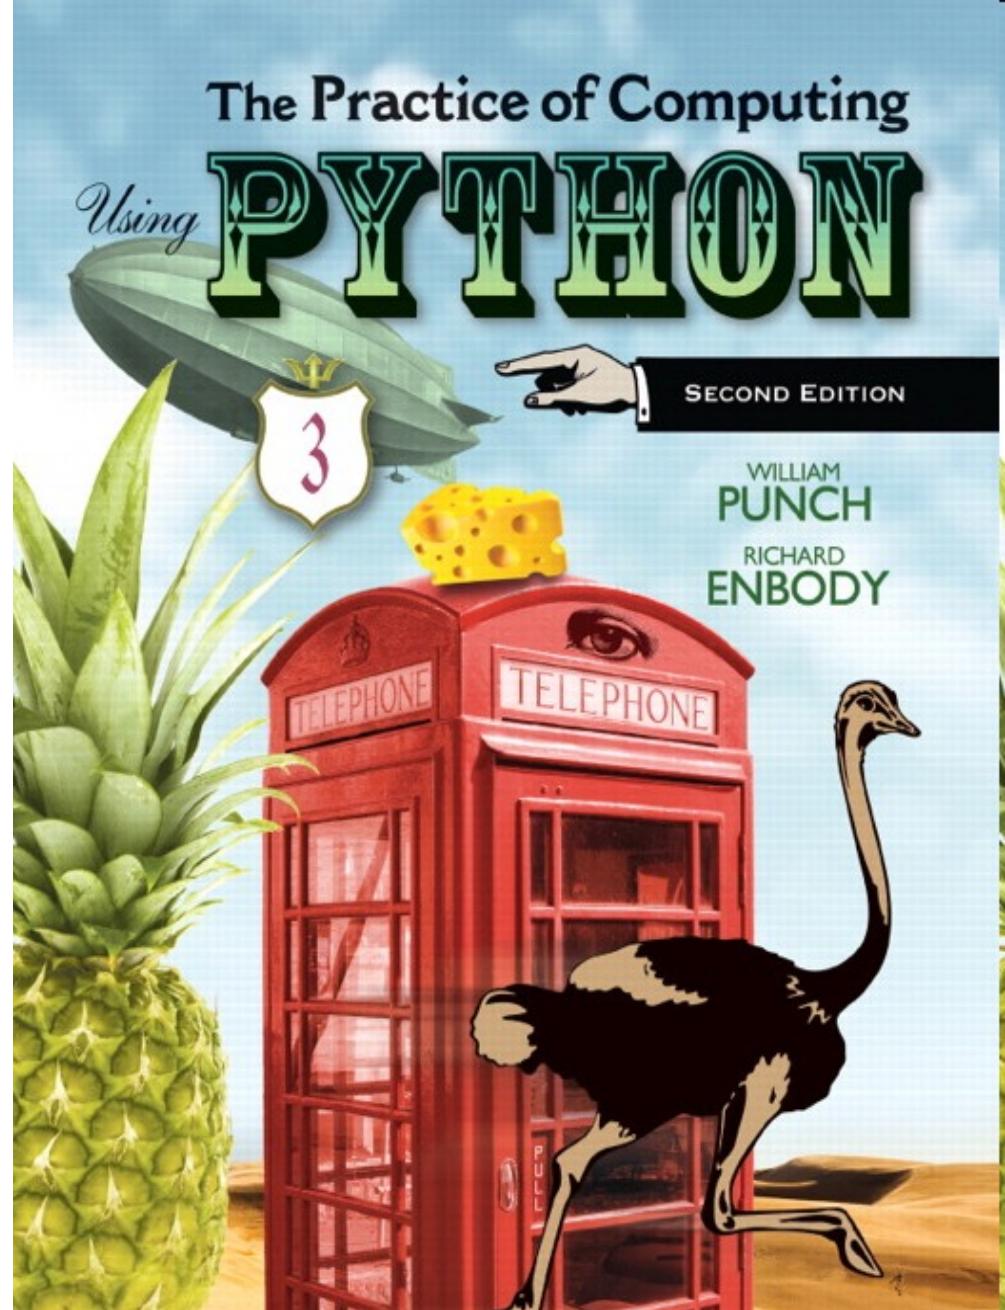


# Chapter 0

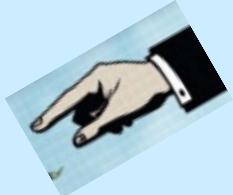
## The Study of Computer Science

PEARSON



ALWAYS LEARNING

# What is Computer Science?

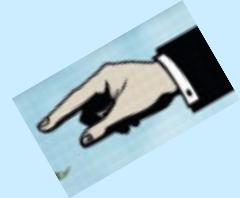


# Definition

- Computer science is a discipline that involves the understanding and design of computers and computational processes.

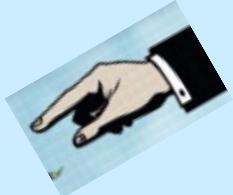


# A well-educated computer scientist should be able to...



- apply the fundamental concepts and techniques of
  - computation,
  - algorithms, and
  - computer design
- to a specific problem





# ...including

- detailing of specifications,
- analysis of the problem,
- provide a design that
  - functions as desired,
  - has satisfactory performance,
  - is reliable and maintainable,
  - and meets desired cost criteria.

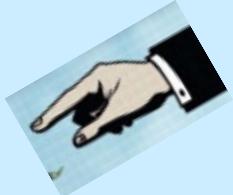




# Aspects of Computer Science

- Theory of computation
- Computational Efficiency
- Algorithms and Data Structures
- Parallel Processing
- Software Engineering
- Others

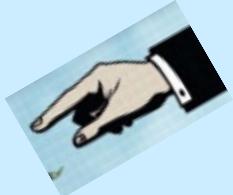




# Our goals

- Our goals are not to just write copious amounts of code, our goals are to:
- increase our problem solving skills
- design good solutions to problems
- test somehow how well they are indeed solutions to the problem
- provide the solution as a readable document

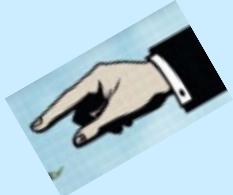




# Why is this hard?

- I cannot precisely explain why it is hard, only that it is indeed hard.
- Typically quote is “Never have I worked so hard and gotten so low a grade”





# An analogy

Let us say that you have signed up to study French poetry (how about Marot) in the original

You have two problems:

- you don't speak French
- you don't know much about poetry



# Clement Marot

## 1496- 1544

- *Ma mignonne*
- *Je vous donne*
- *Le bon jour;*
- *Le séjuour*
- *C'est prison.*
- *Guérison*
- *Recourez,*
- *Puis ouvrez*
- *Votre porte*
- *Et qu'on sorte*
- *Vitement,*
- *Car Clément*
- *Le vous mande.*
- *Va, friane*
- *De ta bouche,*
- *Qui se couche*
- *En danger*
- *Pour mange*
- *Confitures;*
- *Si tu dures*
- *Trop malade,*
- *Couleur fade*
- *Tu Prendras,*
- *Et perdras*
- *L'embonpoint.*
- *Dieu te doint*
- *Santé bonne,*
- *Ma mignonne.*



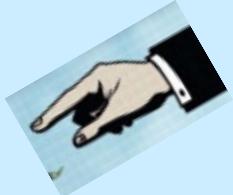
# Crappy- literal translation

- *My sweet/cute [one] (feminine)*
- *I [to] you (respectful) give/bid/convey*
- *The good day (i.e., a hello, i.e., greetings).*
- *The stay/sojourn/visit (i.e., quarantine)*
- *[It] is prison.*
- *Cure/recovery/healing (i.e., [good] health)*
- *Recover (respectful imperative),*
- *[And] then open (respectful imperative)*
- *Your (respectful) door,*
- *And [that one (i.e., you (respectful)) should} go out*
- *Fast[ly]/quick[ly]/rapid[ly],*
- *For/because Clement*
- *It (i.e., thusly) [to] you(respectful) commands/orders.*
- *Go (familiar imperative), fond-one/enjoyer/partaker*
- *Of your (familiar) mouth,*
- *Who/which herself/himself/itself beds (i.e., lies down)*
- *In danger;*
- *For/in-order-to eat*
- *Jams/jellies/confectionery.*
- *If you (familiar) last (i.e., stay/remain)*
- *Too sick/ill,*
- *[A] color pale/faded/dull*
- *You )familiar) will take [on],*
- *And [you (familiar)] will waste/lose*
- *The plumpness/stoutness/portliness (i.e., well-fed look).*
- *[may] God [to] you (familiar) give/grant*
- *Health good,*
- *My sweet/cute [one] (feminine).*



- *My sweet dish,*
- *You I wish*
- *A good day.*
- *Where you stay,*
- *Is a jail.*
- *Though so pale,*
- *Leave your bed,*
- *Regain red.*
- *Ope' your door*
- *Stay not, poor*
- *Child; gain strength*
- *And at length,*
- *Steve does urge,*
- *Please emerge.*
- *Then go eat*
- *Jam so sweet.*
- *Lying ill*
- *Means you will*
- *become too thin -*
- *Merely skin*
- *Cov'ring bone;*
- *Regretted tone.*
- *Eat again,*
- *Avoid the fen.*
- *God grant thee*
- *Be healthy.*
- *This I wish,*
- *My sweet dish.*

# Decent Trans, S.Jamar



# How does this apply

You have two related problems:

- the syntax of French is something you have to learn
- the semantics of poetry is something you have to learn

You have two problems you have to solve at the same time.

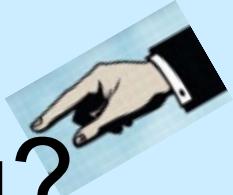


# Programming, Syntax and Semantics



- You have to learn the syntax of a particular programming language
  - many details about the language, how to debug and use it
- You have to learn about problem solving and how to put it down on computer.
- There probably is no better way. It's hard!



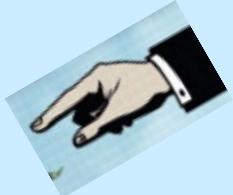


# Computers & problem solving?

This is both the problem and difficulty of computers.

- The promise (perhaps the hope) of computers is that, somehow, we can embed our own thoughts in them. To some extent we can!
- The problem is the difficulty of doing so, and the stringent requirements, the real rigor, required to put simple thoughts into a working program.





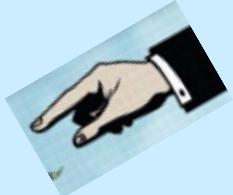
# Focus of Computer Science

There are two foci for computer science

- Learning the difficult task of truly “laying out” a problem-solving task
- Providing tools to make this process as easy (though it will never be “easy”) as possible.

Your focus should be on problem-solving, and adding rigor/focus to your ability to do problem-solving.



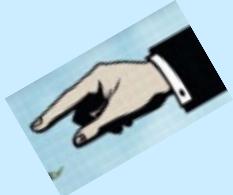


# Good Program (1)

What makes a good program?

- a program is a reflection of the writer and their thoughts
- First, you must have some thoughts! The difficulty for most people is to figure out what has to be done, the problem solving, before writing a program

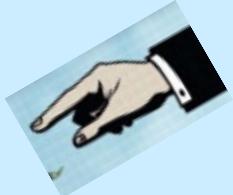




# Rule 1

Think before you program!

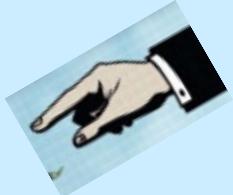




# Good Program (2)

- It will be said repeatedly that the goal of a program is not to run, but to be read.
- A program communicates with other people as well. It stands as a document to be read, repaired and, yes, run

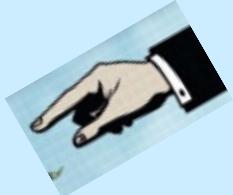




# Rule 2

A program is a human-readable essay on problem solving that also happens to execute on a computer.



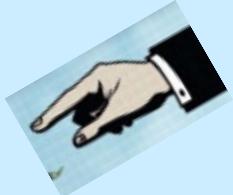


# Why Python?

This book utilizes the programming language known as Python.

## Why?





# Why Python (1): Simpler

- Python is a simpler language
- Simpler means:
  - Fewer alternatives (one way to do it)
  - Better alternatives (easier to accomplish common tasks)
- This allows us to focus less on the language and more on problem solving

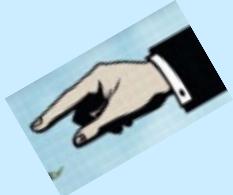


# Why Python(2): Best Practices



- Many of the best parts of other languages are included in Python
  - data structures (lists, dictionaries)
  - control (iteration, exceptions)
  - many packages for common tasks
- Python is often described as "batteries included"

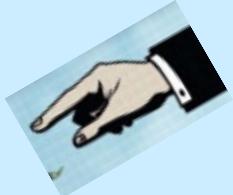




# Why Python(3): User base

- While we want to (and will) teach the fundamentals of computer science, we want what you learn to be useful
- Python is Open Source:
  - freely available
  - large user base constantly contributing
  - new packages available to meet changing needs

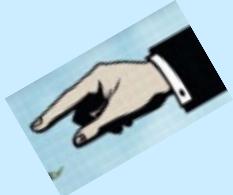




# Why Python (4): Useful

- As a result, Python is more generally useful for getting work done.
- One course in Python makes you a capable (though not expert) programmer
- Can use available packages and your new skill to solve problems.





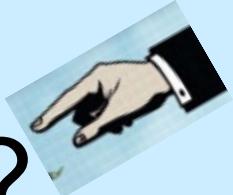
# Computational Thinking

Having finished this course, we want you to have the following thought in your subsequent college career.

“Hey, I’ll just write a program for that”. For us, that is “computational thinking”

Python allows this to happen more readily.



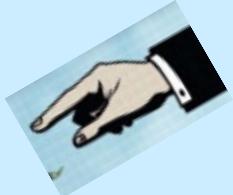


# Is Python the best language?

- The answer is no. This is because there is no best language.
- Computer languages, like tools, are suited for different tasks (what's the best shovel? Depends on what you are doing).
- For introductory students, we think Python is a very good language.



# What is a computer?



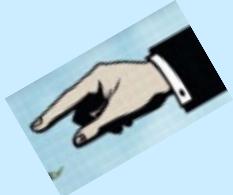
# Computer as a toaster

50 years ago when computers came into common existence, they were expensive, rare items used for research.

Today they are as common as a toaster  
(likely more common, do you own a toaster?)

Good to know a little bit about your toaster

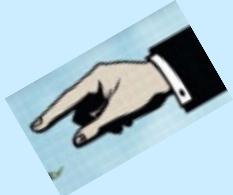




# Computers do computation

- Kind of obvious, but a computer is something that does computation.
- What's interesting is what counts as computation

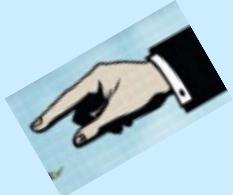




# Human Brain

- In the 1960's, when computers were becoming more prevalent, they were commonly called "electronic brains"
- However, the brain as a computer has very little in common with the modern computer (except that they both do computation).



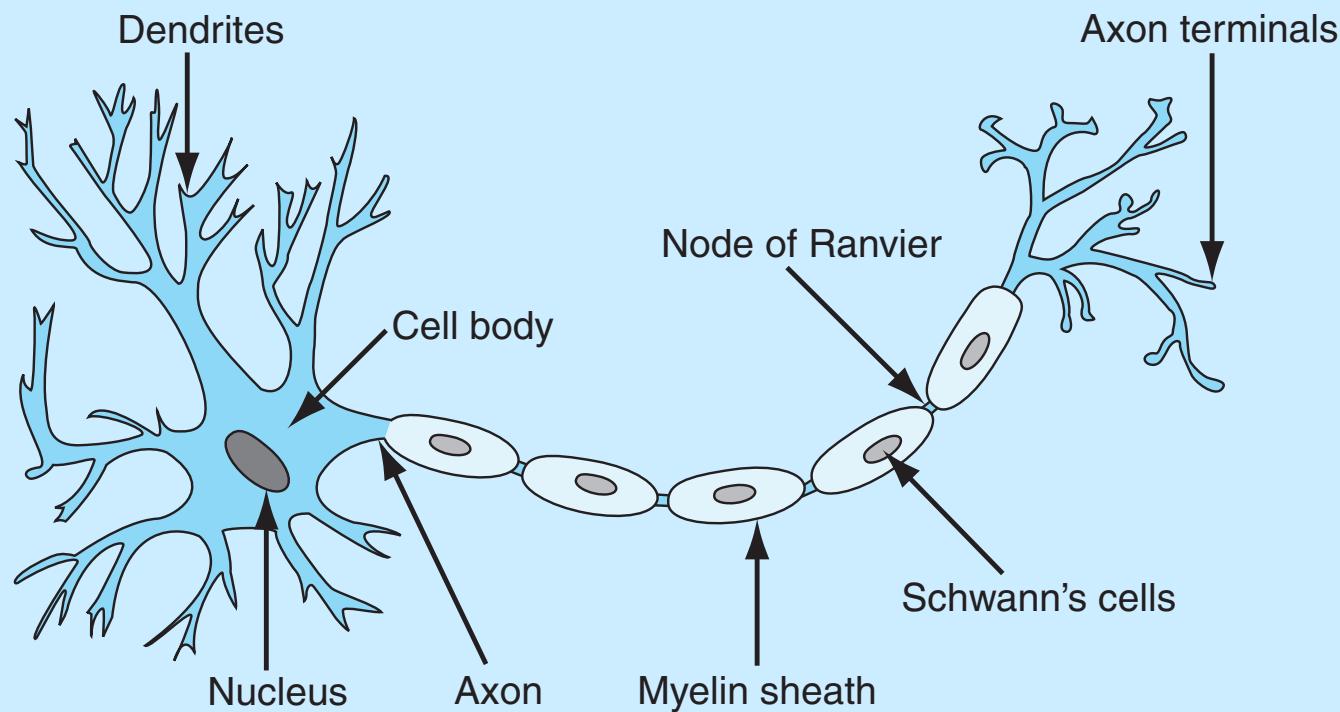


# Neuron

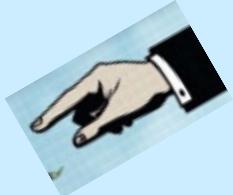
- The human brain consists of on the order of 100 billion ( $10^{11}$ ) small cells called neurons
- Neurons have a simple task, to act as a kind of switch



## Structure of a Typical Neuron



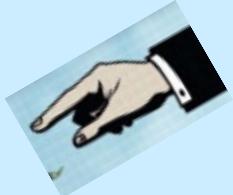
**FIGURE 0.3** An annotated neuron. [SEER Training Modules, U.S. National Institutes of Health, National Cancer Institute]



# Transmission of signal

- A neuron “fires”, sends a signal, when enough chemical transmitters accumulate at its dendrites.
- When it fires, the signal, as a chemical/electrical impulse moves down the axon
- When the signal reaches the terminals, more chemical transmitters are released to other neurons.

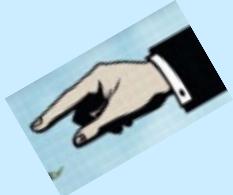




# Interesting facts

- Signal transmission is very slow
  - millions times slower than a computer
- Neuron's require recovery time before they can fire again.
- What causes a firing depends on the number of signals it receives at its dendrites and how fast they arrive.
  - that is, it depends on the number of connections.

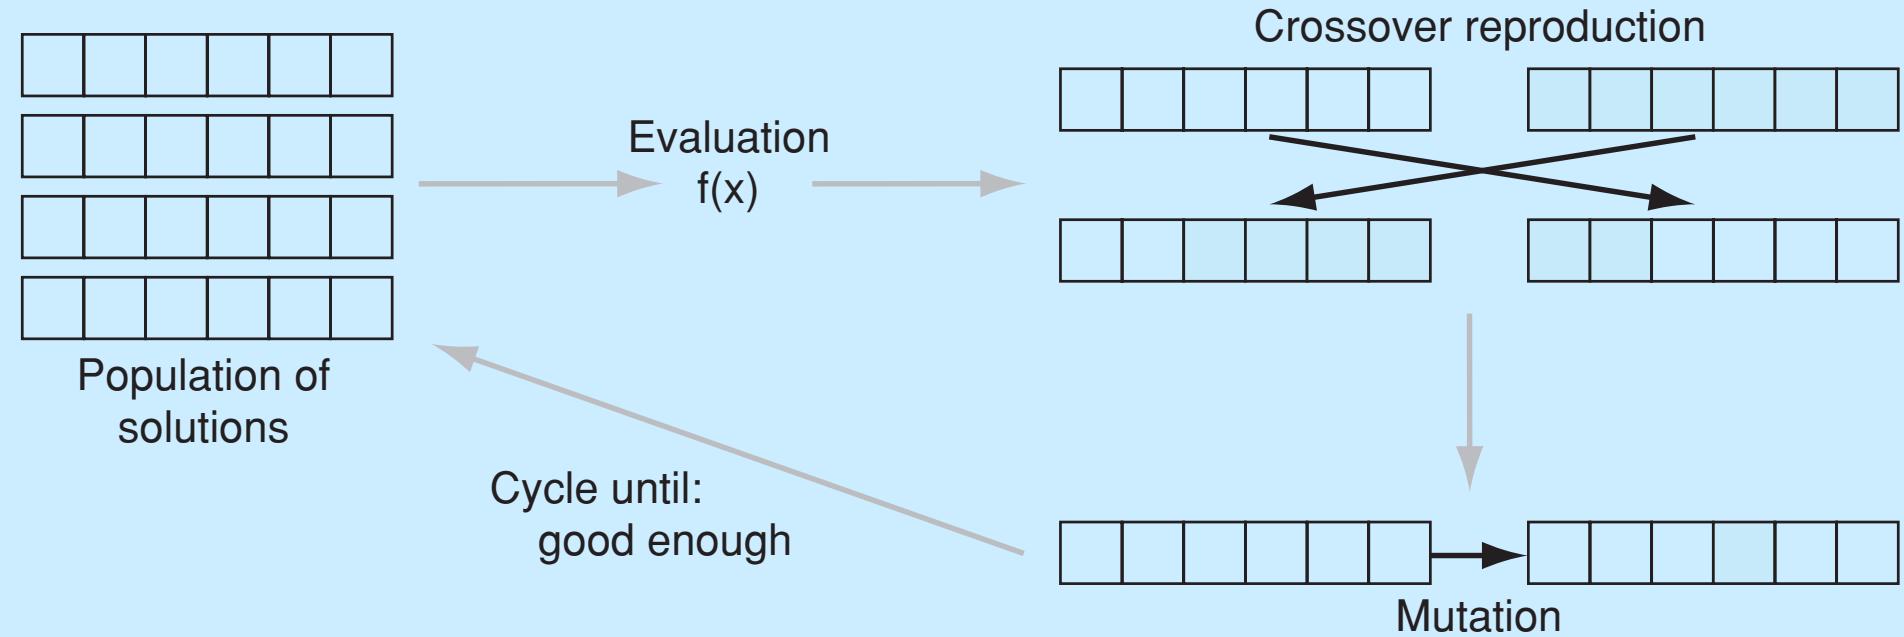




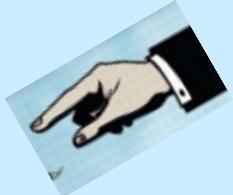
# Evolution as Computation

- There is an entire field devoted to the use of the principles of Darwin and Mendel to solve a problem: evolutionary computation
- Evolution focuses on solving a problem by generating solutions (i.e. organisms) that adapt to the present environment





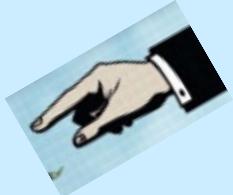
**FIGURE 0.4** A genetic algorithm.



# Facts

- maintains a population of solutions
  - solutions are an encoding of the problem
- Each solution is evaluated
- Good solutions tend to be used to create new solutions, variations on themselves
- Over time, the population contains better and better solutions





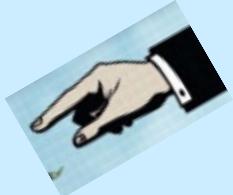
# People!

Computers used to refer to *people*

In WWII, computers were people who did difficult calculations *by hand*, for things like ballistic tables.



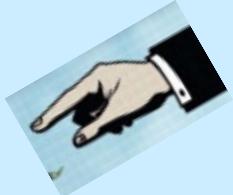
# The Modern Computer



# It's all about the switch

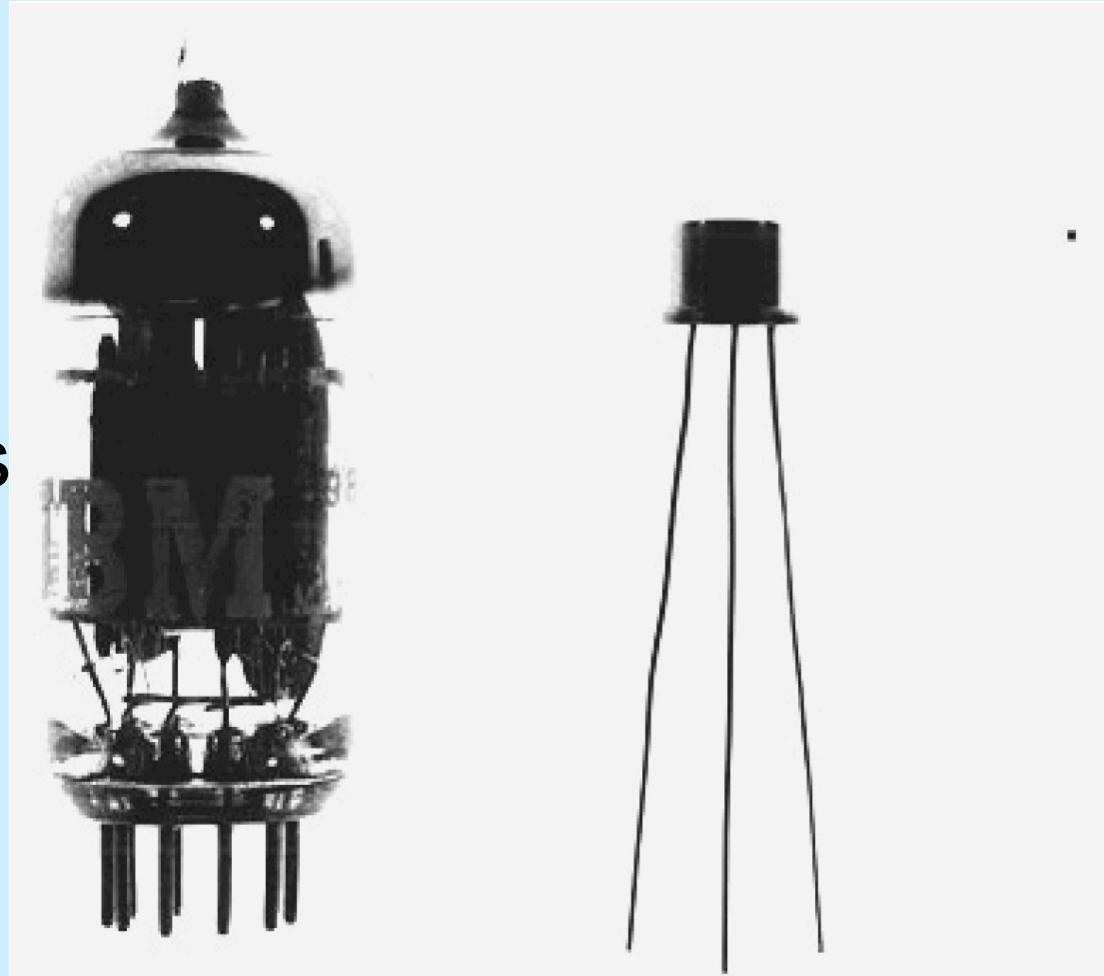
- The basic component of most digital circuitry is nothing more complicated than a simple switch.
- A switch's function is pretty obvious, said in a number of different ways
  - On or Off
  - True or False
  - 1 or 0

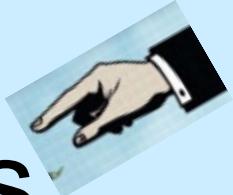




# Electronic switch

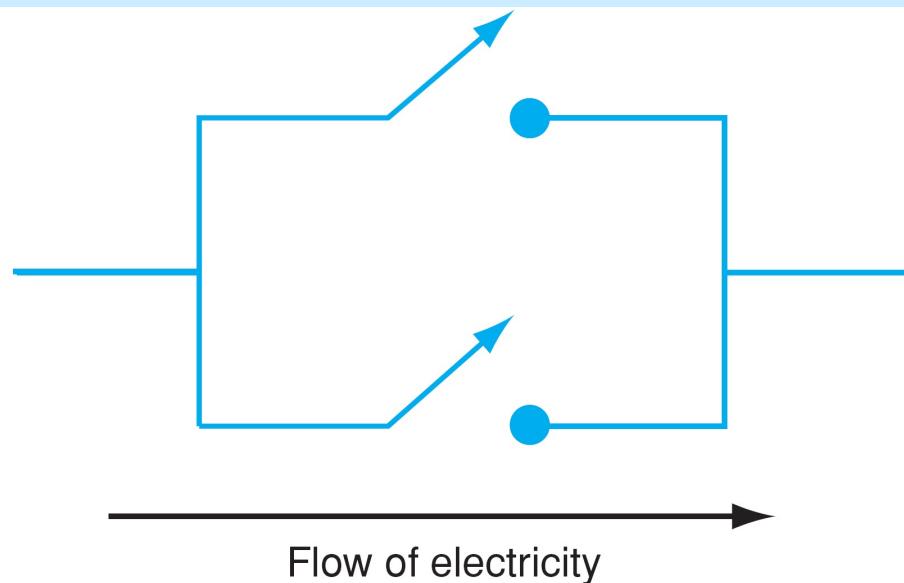
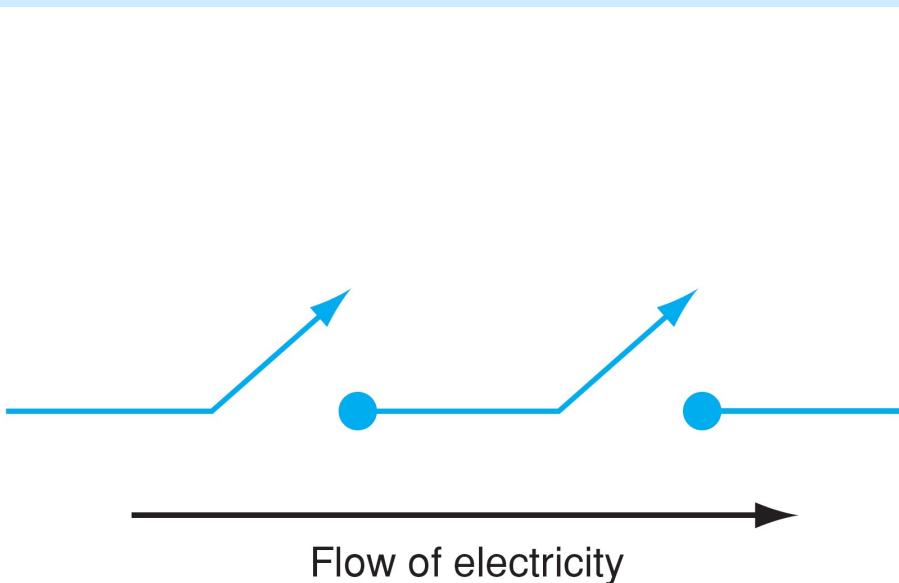
- Early computers used vacuum tubes as switches
- Later, transistors were used as substitutes



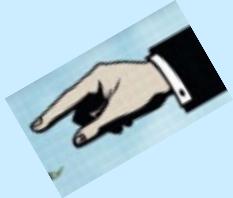


# Switches for Boolean Circuits

- Switches can be used to construct more complicated functions, such as Boolean circuits (and on left, or on right)

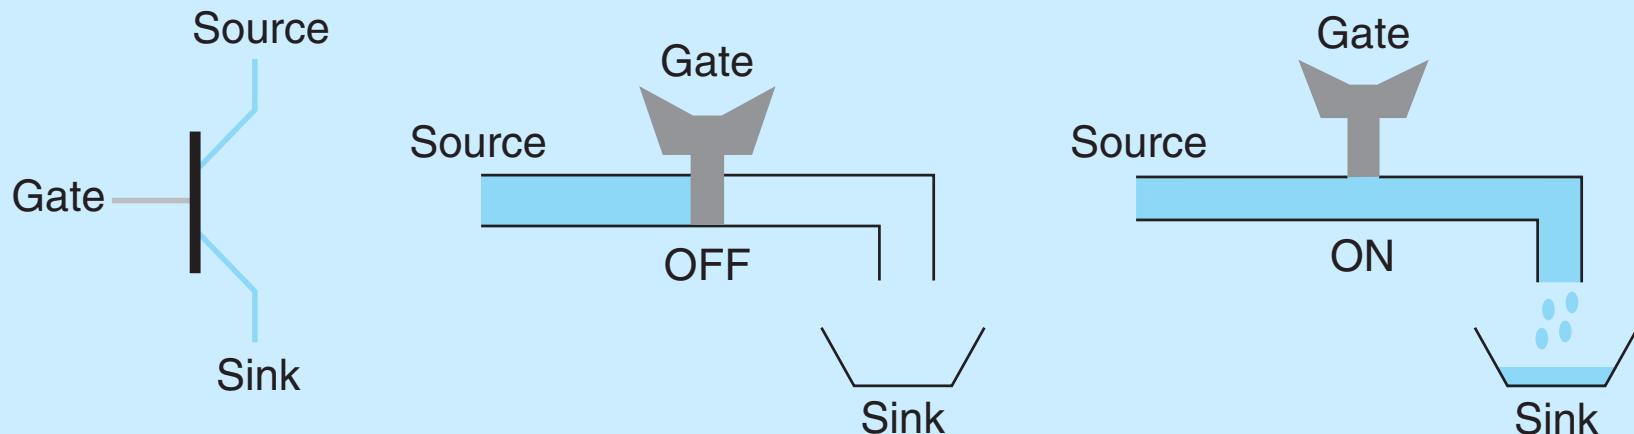


# Transistor is the key



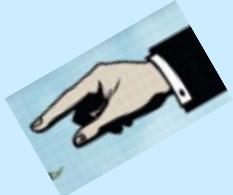
Transistor is an electronic switch:

- if gate has current, then signal flows from source to sink
- if gate has no current, no flow



**FIGURE 0.8** A diagram of a transistor and its equivalent “faucet” view.



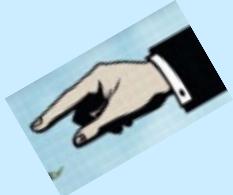


# But a very interesting switch

Transistors have three interesting “features” that have made them the fundamental element of the computer revolution

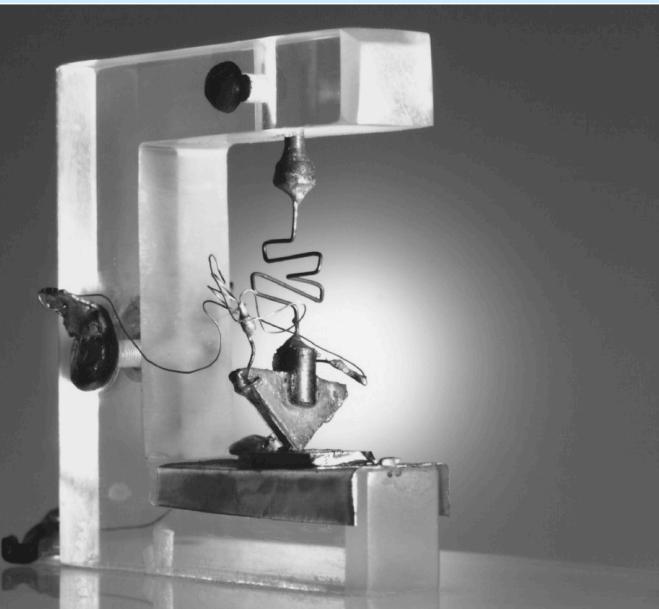
- size
- quantity
- speed



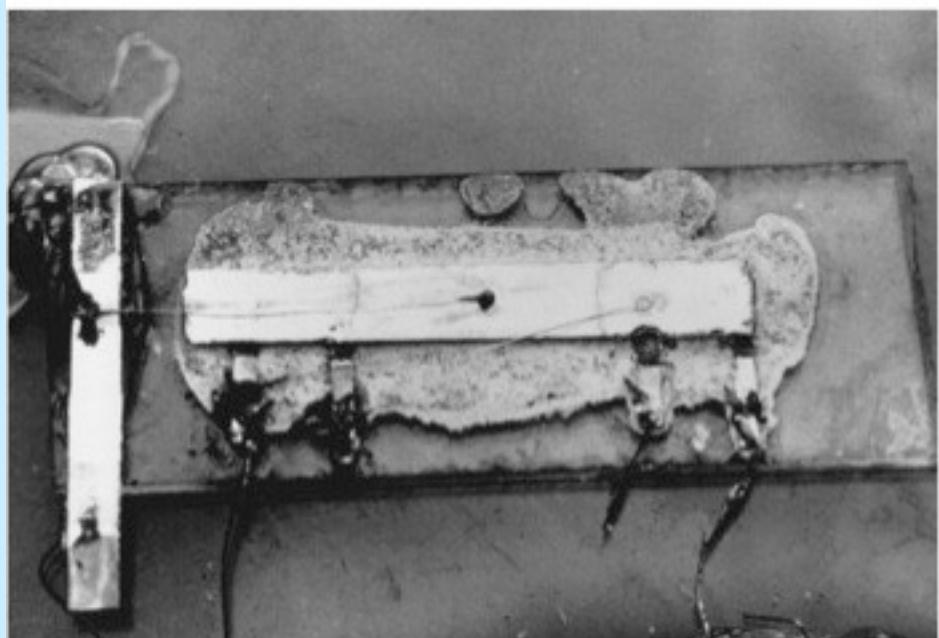


# Size

Originally very large

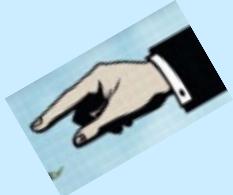


Shockley transistor



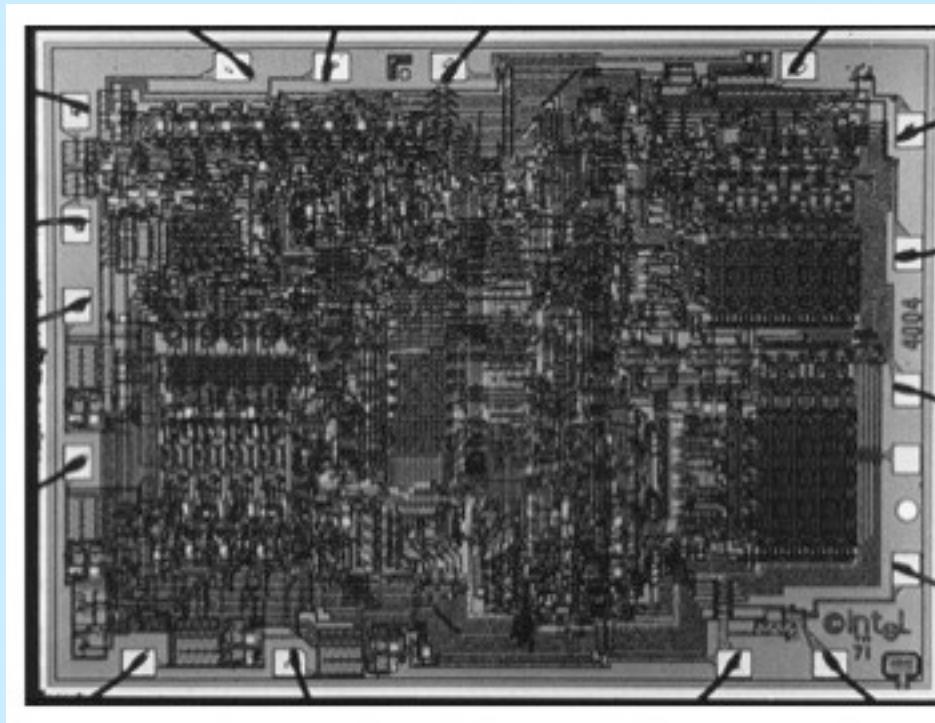
Kilby integrated circuit

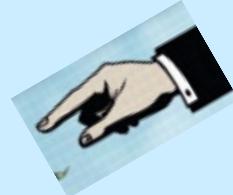




# Intel's first cpu

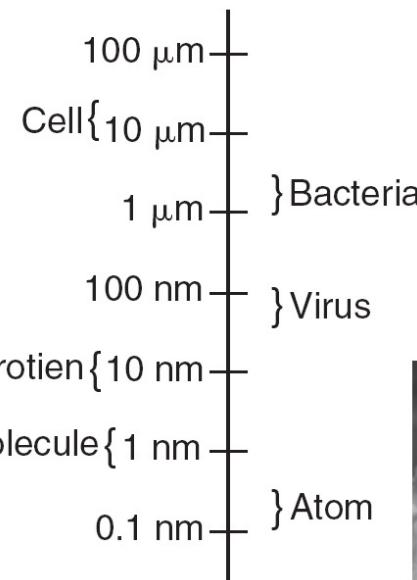
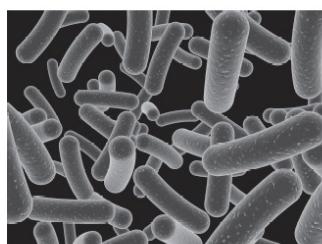
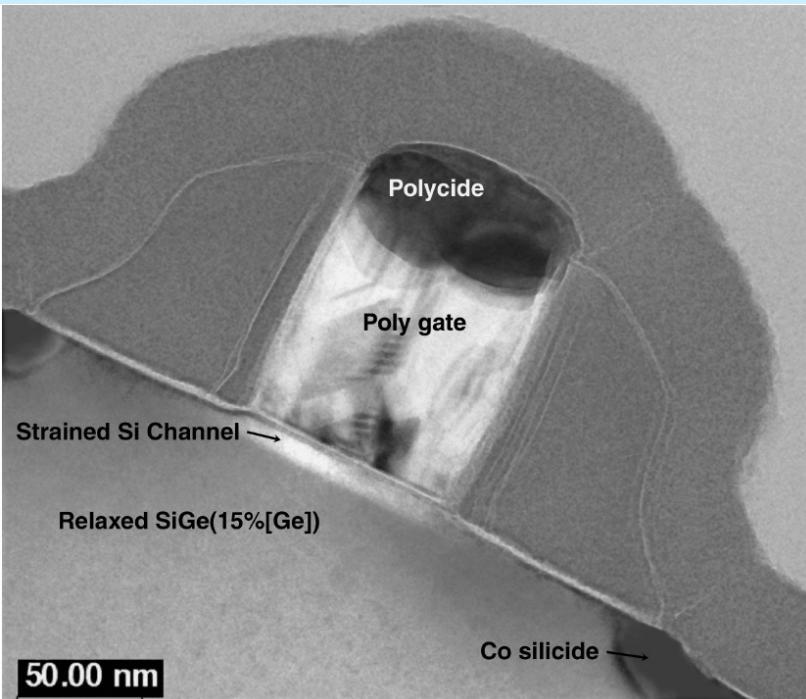
By 1971, Intel had created a “computer on a chip”, the 4004 microprocessor, the size of a fingernail with 2300 transistors



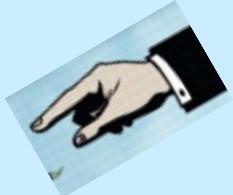


# But now, *really* small!

Chips now have gates measured in billionths of a meter (nanometers, nm)



# Smaller means more



Year	Transistor Count	Model
1971	2,300	4004
1978	29,000	8086
1982	134,000	80286
1986	275,000	80386
1989	1,200,000	80486
1993	3,100,000	Pentium
1999	9,500,000	Pentium III
2001	42,000,000	Pentium 4
2007	582,000,000	Core 2 Quad
2011	2,600,000,000	10-core Westmere

**TABLE 0.1** Transistor counts in Intel microprocessors, by year.

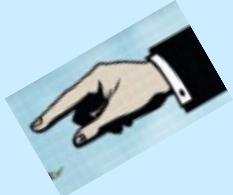




# GPUs too!

<b>Processor</b>	<b>count</b>	<b>Date</b>	<b>Co.</b>	<b>Process</b>	<b>Area</b>
G80	681,000,000	2006	NVIDIA	90 nm	480 mm <sup>2</sup>
RV770	956,000,000	2008	AMD	55nm	
GT200	1,400,000,000	2008	NVIDIA	55 nm	576 mm <sup>2</sup>
HD5800	2,154,000,000	2009	AMD	40 nm	334 mm <sup>2</sup>
GF100	2,900,000,000	(future)	NVIDIA	40 nm	



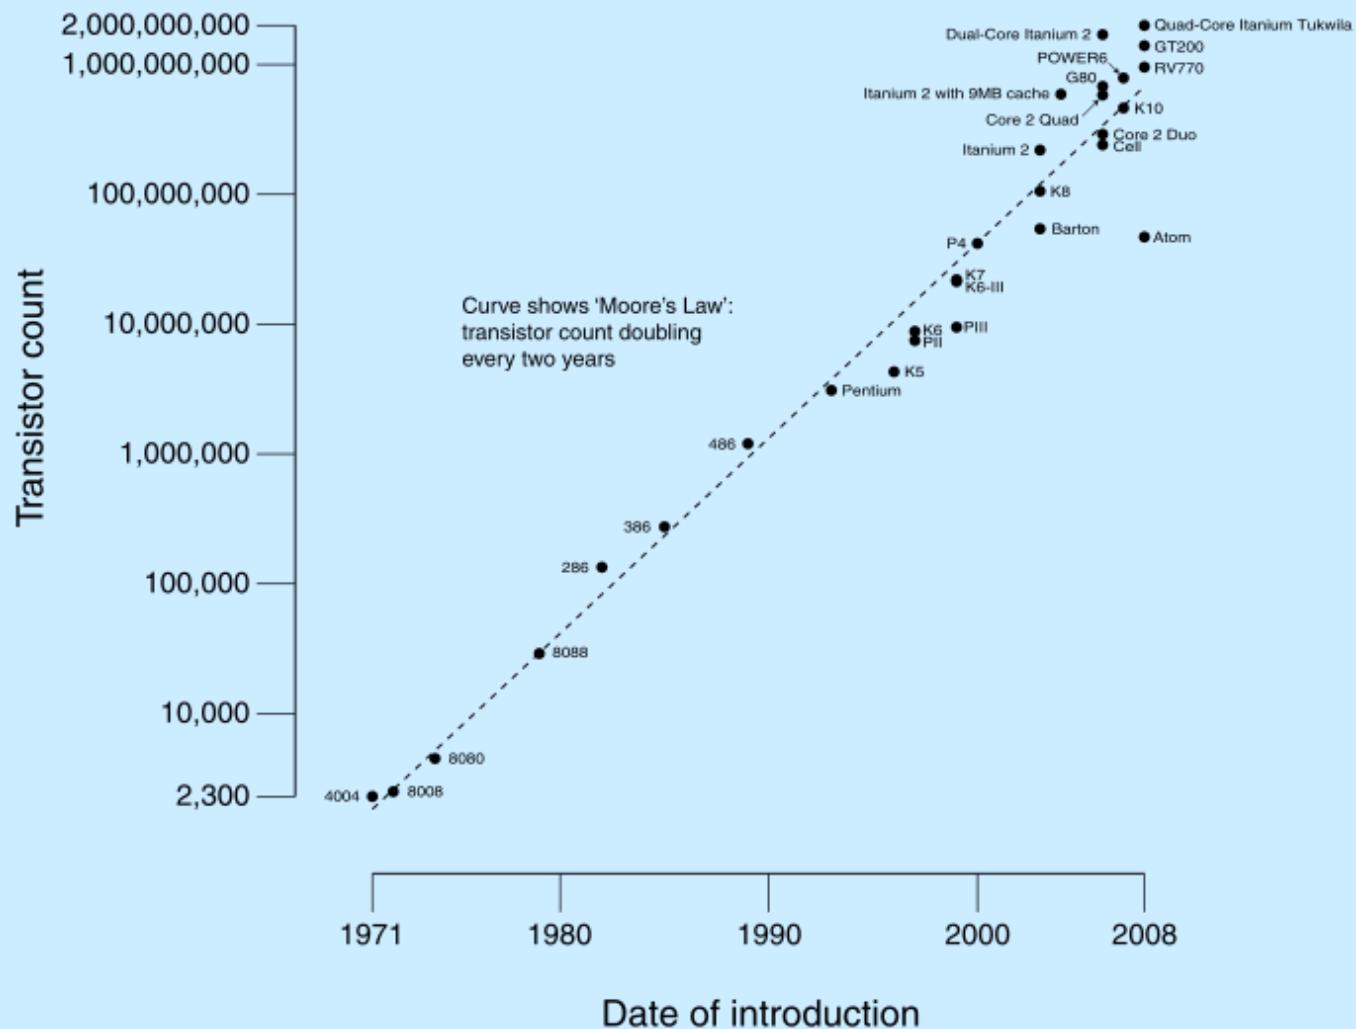
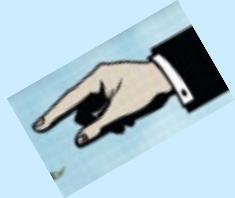


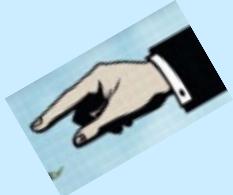
# Moore's Law

- Gordon Moore is one of the founders of the chip maker Intel
- in 1965, he has observed (over that last 15 years or so) the growth rate of the number of transistors in a circuit
- Made a famous prediction



# CPU Transistor Counts 1971-2008 & Moore's Law



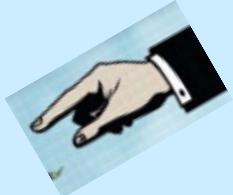


# The “law”

“The complexity for minimum component costs has increased at a rate of roughly a factor of two per year ... Certainly over the short term this rate can be expected to continue, if not to increase. Over the longer term, the rate of increase is a bit more uncertain, although there is no reason to believe it will not remain nearly constant for at least 10 years. That means by 1975, the number of components per integrated circuit for minimum cost will be 65,000. I believe that such a large circuit can be built on a single wafer”

Electronics Magazine 19 April 1965

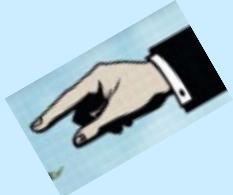




# What it means

- Roughly, since 1965, the number of transistors on a chip doubles every 18 months for approximately the same cost
- Often quoted as the speed of a cpu doubling every 18 months for the same cost
- speed and density were related





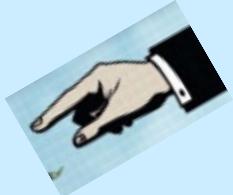
# How fast

Typical processor this day runs at GHz speeds. That is 1 billion operations a second.

Year	CPU	Instructions/second	Clock Speed
1971	Intel 4004	1 MIPS	740 kHz
1972	IBM System/370	1 MIPS	?
1977	Motorola 68000	1 MIPS	8 MHz
1982	Intel 286	3 MIPS	12 MHz
1984	Motorola 68020	4 MIPS	20 MHz
1992	Intel 486DX	54 MIPS	66 MHz
1994	PowerPC 600s (G2)	35 MIPS	33 MHz
1996	Intel Pentium Pro	541 MIPS	200 MHz
1997	PowerPC G3	525 MIPS	233 MHz
2002	AMD Athlon XP 2400+	5,935 MIPS	2.0 GHz
2003	Pentium 4	9,726 MIPS	3.2 GHz
2005	Xbox360	19,200 MIPS	3.2 GHz
2006	PS3 Cell BE	10,240 MIPS	3.2 GHz
2006	AMD Athlon FX-60	18,938 MIPS	2.6 GHz
2007	Intel Core 2 QX9770	59,455 MIPS	3.2 GHz
2011	Intel Core i7 990x	159,000 MIPS	3.46 GHz

## How fast is a nanosecond?





# Nanoseconds

- Speed of light = 186,000 miles/sec,  
 $3 \times 10^8$  meters/sec (actually 299,792,458)
- nanosecond =  $10^{-9}$  seconds
- thus 30 centimeters or 11.8 inches

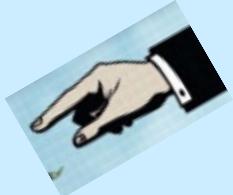




# speed pushing up against physics

- the “clock” is like a drummer in the band. The faster it beats, the faster the operations.
- in 1 clock tick (at 1 GHz), electricity can travel 1 ft. At 2GHz, 6 inches. At 4GHz, 3 inches
- It becomes very difficult to get electricity to travel any distance in that short of time!

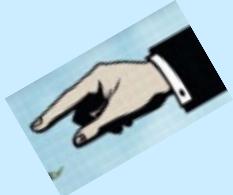




# How to get around physics

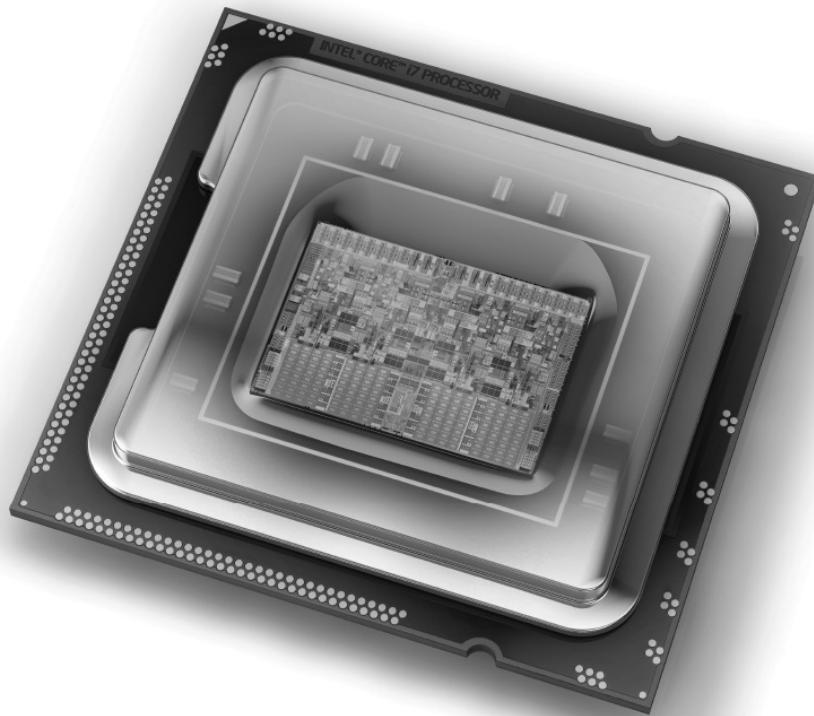
- So if physics is getting in the way (and it is), you find a way to get around it.
- If you can't make processors faster, what do you do?



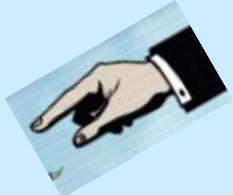


# Not faster, but more CPUs

Now more transistors mean more CPUs on a single chip. Below a quad-core CPU



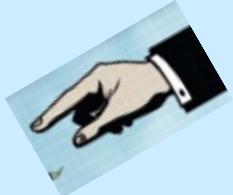
# Computer Parts



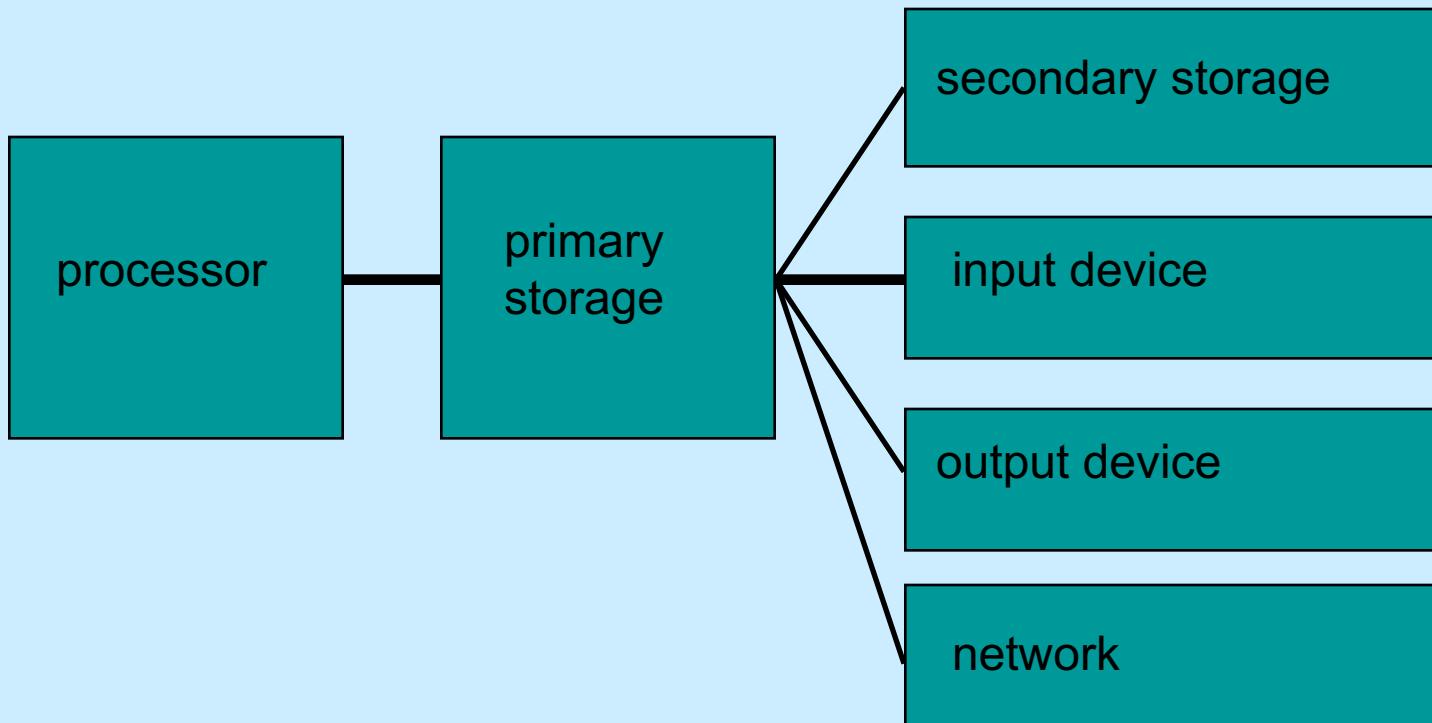
# Main Components

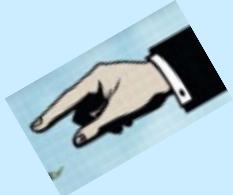
- People!!!
- Hardware
  - Physical Devices: processor, memory, keyboard, monitor, mouse, etc.
- Software
  - Executable Programs: word processor, spread sheet, internet browser, etc.





# Hardware

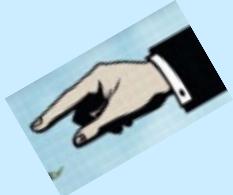




# Processor

- The processor is the “brain” of a computer.
- The processor controls the other devices as well as performing calculations

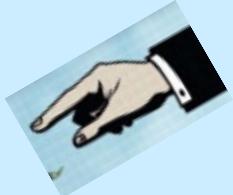




# Primary Storage

- stores instructions and data for current program(s)
- other names: primary or main memory, RAM (Random Access Memory)
- memory is dynamic so it requires power to retain information
- often hundreds of Megabytes (million-bytes)

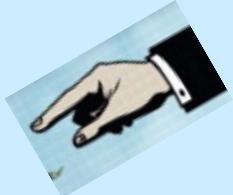




# Peripheral Devices

- Secondary storage devices
  - disk (hard & floppy), tape, usb drives, flash drives, etc.
- Input devices
  - keyboard, mouse, camera, mic, etc.
- Output devices
  - monitor, printer, speaker, etc.
- Network
  - wireless, bluetooth, ethernet, etc.

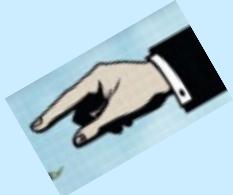




# Secondary Storage

- nonvolatile -- information is recorded magnetically so power is not needed
- disks hold Gigabytes (billions of bytes)
- cheap, but slow
  - RAM access is a hundred CPU clock ticks
  - disk access is a million CPU clock ticks
- not directly accessed by CPU

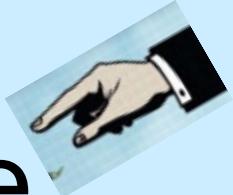




# Software

- the programs available for execution
- simple classification
  - system software
  - application software



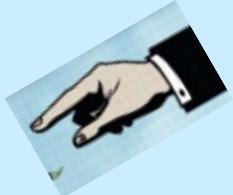


# Fetch-Decode-Execute-Store

CPU operates on a simple (but fast) cycle:

- **Fetch**: fetch instruction from memory
- **Decode**: Decode requirements (args, etc.)
- **Execute**: Perform operation
- **Store**: move needed results to memory
- Repeat

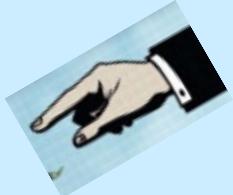




# What is a program?

- A sequence of instructions written in machine language that tells the CPU to take action (i.e. add two numbers) in a specific order
- In this course we will learn to create programs

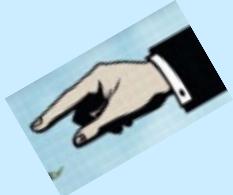




# Program Storage

- machine language instructions are encoded as bit patterns (binary, remember our transistors)
- memory can only hold binary info.
- a bit is represented by two-states, e.g. L-R magnetism, high-low voltage
- it takes many bits to represent reasonable amounts of information

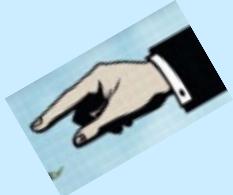




# Binary numbers

- The switch nature of transistors make storing numbers in binary a natural fit.
- Binary is a change of base for our number system, base 2
- In a number, its position represents powers of 2





# Example

- Decimal uses digits 0-9 and positions in a number as powers of 10
  - $735_{10} = 7*10^2 + 3*10^1 + 5*10^0$
- Binary users digits 0,1 and positions in a number as powers of 2
  - $101_2 = 1*2^2 + 0*2^1 + 1*2^0$
- We can convert back and forth
  - $101_2 = 5_{10}$

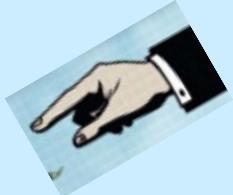




# Rule of thumb: binary to decimal

- The rule: Every 10 powers of 2 is about 3 powers of 10
- $2^{10} = 1024 \approx 10^3 = 1000$  (kilobyte)
- $2^{20} = 1,048,576 \approx 10^6 = 1,000,000$  (megabyte)
- $2^{30} = 1,073,741,824 \approx 10^9 = 1,000,000,000$  (gigabyte)

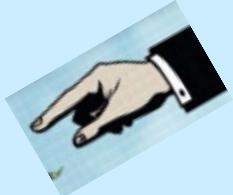




# Bits, Bytes and Words

- Bit: a single 1,0
- Byte: 8 bits
- Word: A computer's standard unit of storage. A typical computer today is a 32 bit computer, meaning its word size is 32 bits (4 bytes)
  - many newer processors are 64 bit!

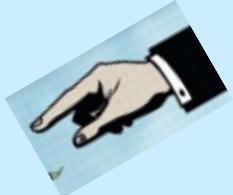




# Floating point

- The approximation of decimal point values such as 1.23, 0.01, 3.14159.
- What is the exact value of  $1/3$  as a decimal point number?
  - It is an infinite sequence which we approximate when using a decimal number
- Remember that floating point is approximate!

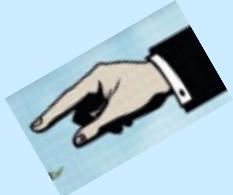




# Representing data in binary

- Letters are represented as numbers!
- Both ASCII and Unicode are encodings for particular letters.
- ASCII (American Standard Code for Information Interchange) was used for English letters. It is superseded by a subset of Unicode called UTF-8 (see appendix E)

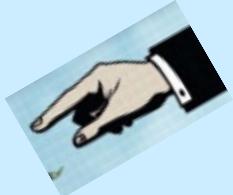




# UTF-8 table (first few rows)

Char	Dec	Char	Dec	Char	Dec	Char	Dec
NUL	0	SP	32	®	64	`	96
SOH	1	!	33	A	65	a	97
STX	2	"	34	B	66	b	98
ETX	3	#	35	C	67	c	99
EOT	4	\$	36	D	68	d	100
ENQ	5	%	37	E	69	e	101
ACK	6	&	38	F	70	f	102
BEL	7	'	39	G	71	g	103
BS	8	(	40	H	72	h	104

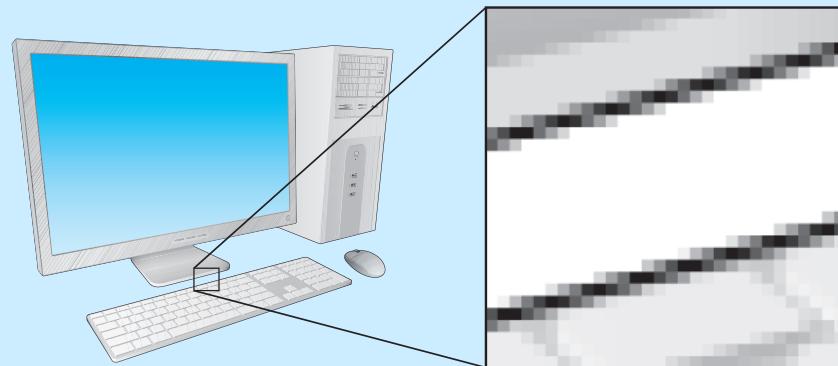




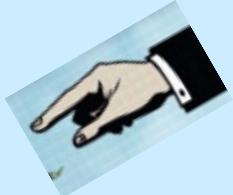
# Images

- Digital images consist of individual colors in a matrix.
- Each individual color is called a pixel.
- The color of a pixel can be encoded using numbers.





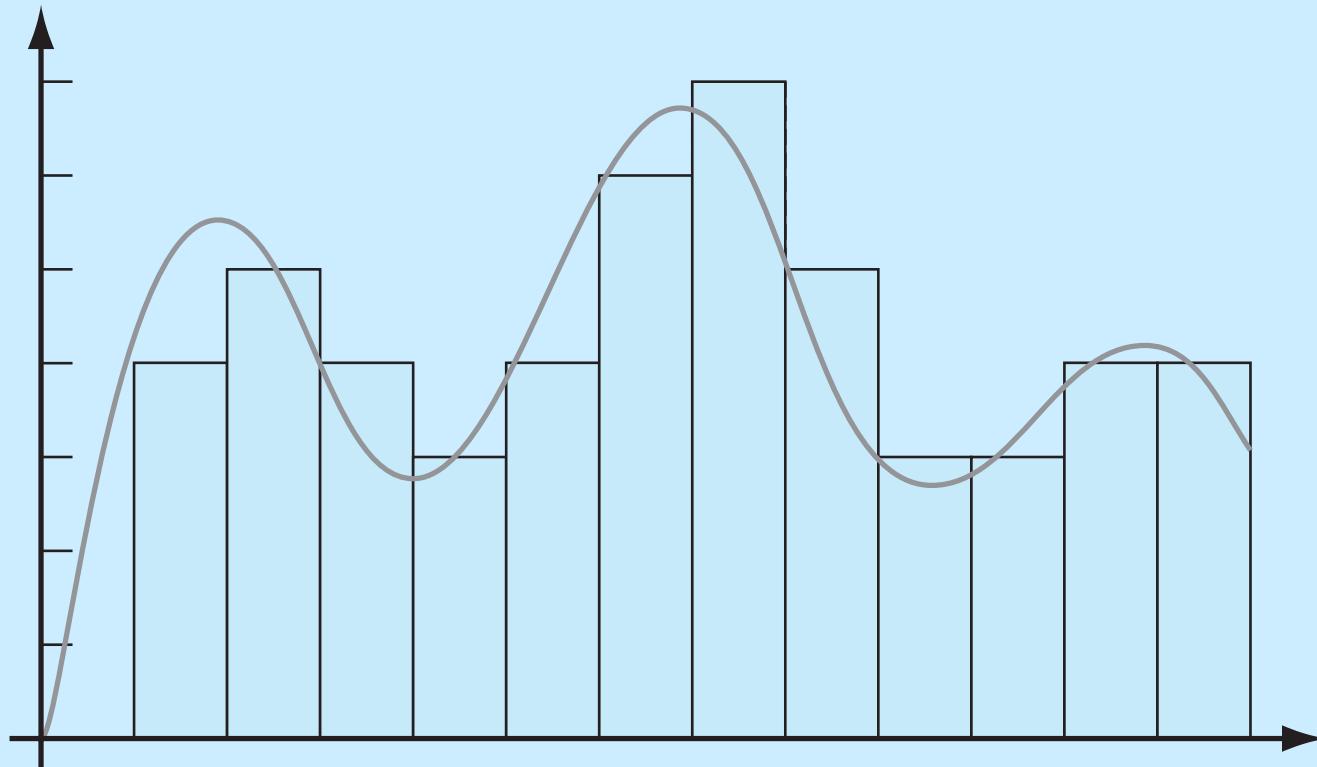
**FIGURE 0.14** A computer display picture with a close-up of the individual pixels. [Juliengrondin/Shutterstock]



# Sound

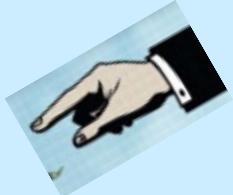
- We can sample a sound wave at a very high frequency and measure the height of the sound.
- We can represent those measurements as numbers and use them to recreate a sound wave





**FIGURE 0.15** A sound wave and the samples of that sound wave (blue bar height) over time.

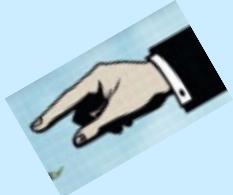
# How many numbers



# Binary

- byte = 8 bits  
e.g. 10110010
- storage devices come in large quantities
  - 1K (kilobyte)=  $2^{10}$  bytes = 1024 bytes
  - 1M (megabyte)=  $2^{20}$  bytes = 1,048,576 bytes
  - 1G (gigabyte)=  $2^{30}$  bytes=1,073,741,824 bytes

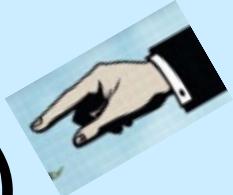




# Kilobyte (1000 bytes)

- 2 Kilobytes: A Typewritten page
- 10 Kilobytes: An encyclopedic page
- 50 Kilobytes: A compressed document
- 100 Kilobytes: A low-resolution photo

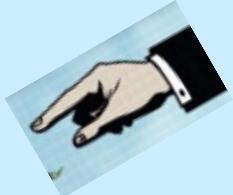




# Megabyte (1 000 000 bytes)

- 1 MByte: A small novel **or** 3.5 inch floppy disk
- 2 MByte: A high resolution photograph
- 5 Mbyte: The complete works of Shakespeare  
**or** 30 seconds of video
- 100 Mbyte: 1 meter of shelved books
- 500 Mbyte: A CD-ROM

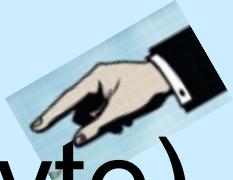




# Gigabyte (1 000 000 000 bytes)

- 1 GByte: A pickup truck filled with paper  
or A symphony  
or A movie
- 20 GBytes: the works of Beethoven as sound
- 50 GBytes: A library floor of books
- 100 Gbytes: standard Blu-ray capacity

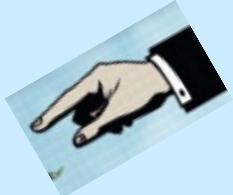




# Terabyte (1 000 000 000 000 byte)

- 1 TByte: All X-rays in a large hospital  
or 50000 trees made into paper and printed  
or Daily rate of EOS data (1998)
- 2 TBytes: An academic research library
- 20 Tb, photos/month on facebook
- 100 TBytes: The US Library of Congress, all types of data, 2009
- 530 Tb, all videos on youtube





# Terabyte In Your Hand

- Terabyte = 1000 Gigabytes  
= 1500 music CD's  
= 200 DVD movies (16 days worth)  
= 1/10 of the Library of Congress
- Western Digital Caviar: 1 Terabyte, \$269 (2006)
  - today, cost is \$109
  - *Seagate Barracuda, 1.5 Tb, \$129*
- History

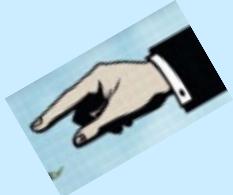
*First disk by IBM in 1956 (4 MB on 50 24-inch platters)*

*1GB in 35 years*

*500 GB in 14 more years*

*1000 GB = 1 TB in 2 more years*





# Petabyte = 1000 Terabytes\* = 1,000,000 Gigabytes

- EMC Symmetrix DMX-3  
1.2 PB = 2400 500GB drives (\$4 million 2006)
- Backblaze (build your own), \$120,000, 2009
- Kazza shared 54 Petabytes (2005)
- Google cluster has 4 Petabytes of RAM,  
processes 20 Pb a day
- All data recorded in 2003: 2500 Petabytes
- WoW, 1.3 Pb to maintain its game

\*(1,125,899,906,842,624 bytes =  $2^{50}$ )

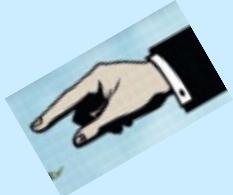


# Petabyte?



- In round numbers, a book is a megabyte. If you read one book a day for every day of your life for 80 years, your personal library will amount to less than 30 gigabytes. Remember a petabyte is 1 million gigabytes so you will still have 999,970 gigabytes left over.
- How many pictures can a person look at in a lifetime (4 Mbytes each)? I can only guess, but 100 images a day certainly ought to be enough for a family album. After 80 years, that collection of snapshots would add up to 30 terabytes. So your petabyte disk will have 970,000 gigabytes left after a lifetime of high quality photos and books.
- What about music? MP3 audio files run a megabyte a minute, more or less. At that rate, a lifetime of listening--24 hours a day, 7 days a week for 80 years--would consume 42 terabytes of disk space. So with all your music and pictures and books for a lifetime you will have 928,000 gigabytes free on your disk.
- The one kind of content that might possibly overflow a petabyte disk is video. In the format used on DVDs, the data rate is about two gigabytes per hour. Thus the petabyte disk will hold some 500,000 hours worth of movies; if you want to watch them all day and all night without a break for popcorn, they will actually fill up your petabyte drive if you have a lifetime of video on it as it will give you 57 years of video





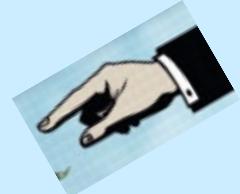
# Exabyte

(1 000 000 000 000 000 000  
bytes)

- 5 Exabytes: All words ever spoken (as text) by human beings.
- 5-8 Eb, global internet traffic monthly
- 500 Eb, total world digital content (2009)



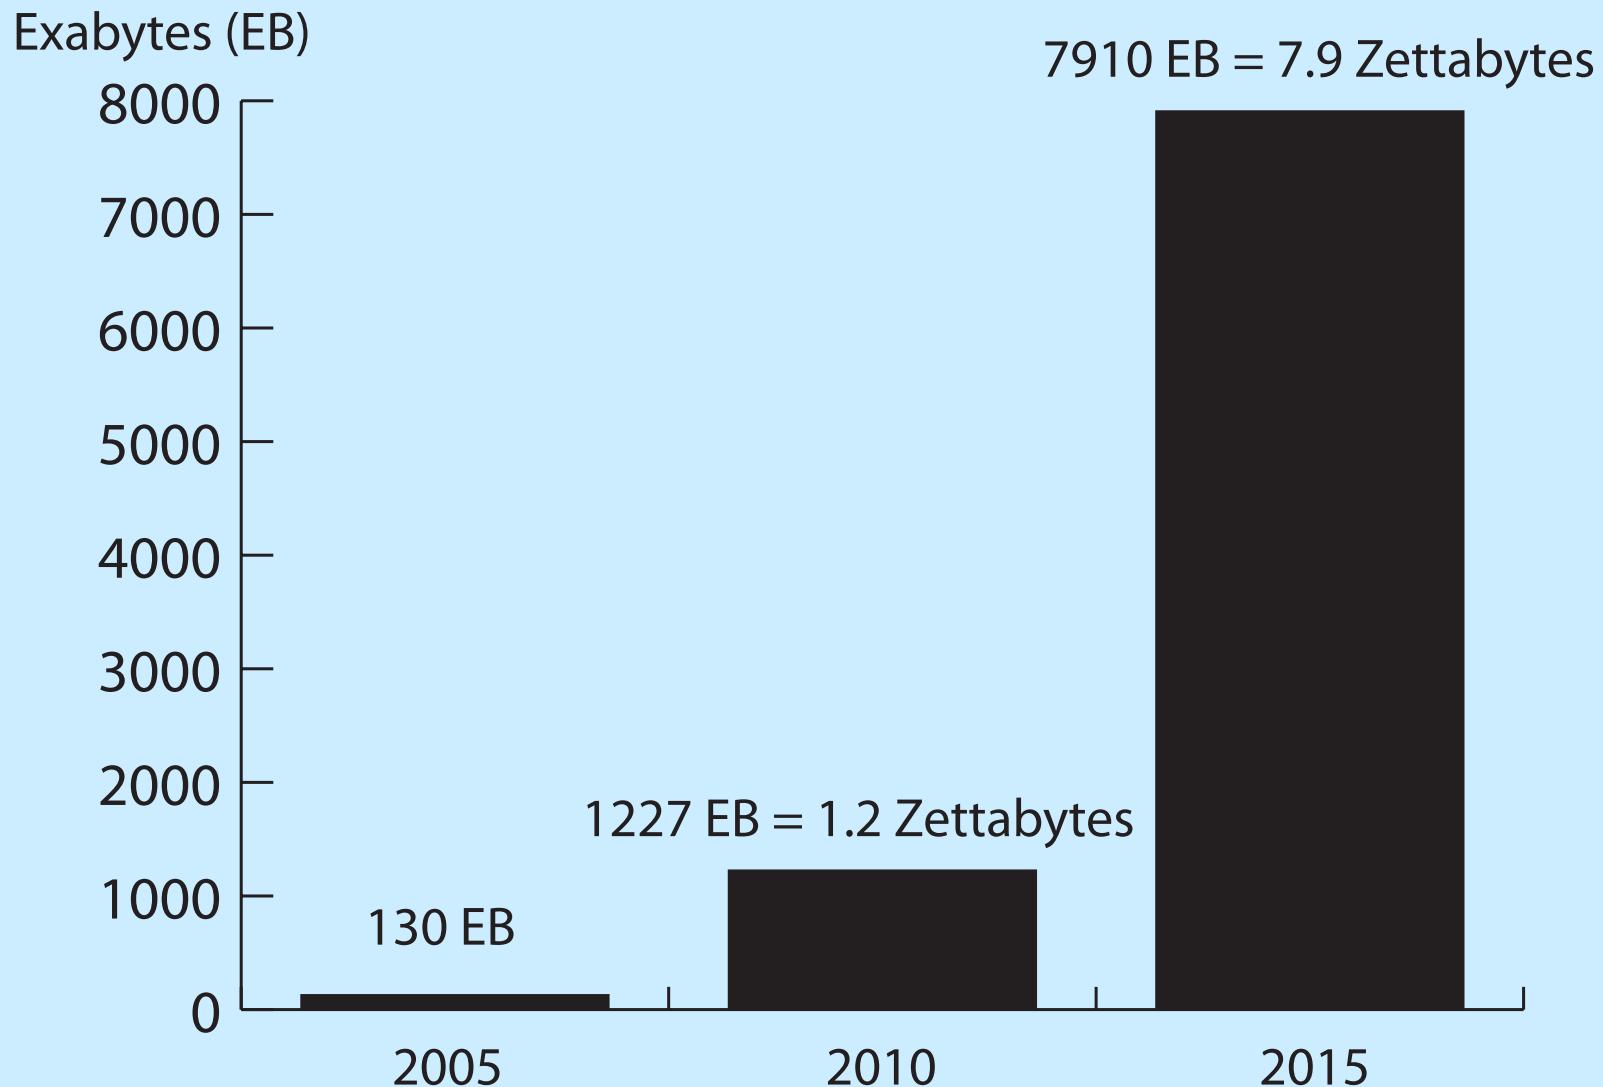
# Zettabyte ( $10^{21}$ )



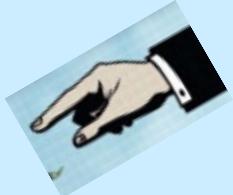
1,000,000,000,000,000,000

- Audio of everything ever spoken, by anyone, at anytime in history, 16KHz, 16 bit audio: 42 Zb





**FIGURE 0.16** Estimates of data created and stored. (IDC Digital Universe)



# The Rules

1. Think before you program
2. A program is a human-readable essay on problem solving that also happens to execute on a computer.

